

Policy Threshold Calibration




Our Current Situation

The results so far are not comforting

...But it's worth seeing what is going on over time

```
In [2]: stop = 1400  
util.plot_rul(tr_pred[:stop], tr['rul'][:stop], figsize=figsize)
```



...And we get the same shapes also on the validation and test set



...And How to Go Forward

Our goal is **not** to regress RUL values with high accuracy

...But rather to define a maintenance **policy** in the form:

$$f(x, w) < \theta \Rightarrow \text{trigger maintenance}$$

- For this, we just need to **stop at the right time**

Our model...

Makes large estimation errors when the RUL is high

- ...But we do not care about those!

Works reasonably well for low RUL values

- ...I.e. exactly where it matters



Threshold Calibration as an Optimization Problem

Given a RUL estimator

...We can choose when to trigger maintenance by calibrating θ

- This is in fact an(other) **optimization problem**
- ...And to formulate it we need a **cost function**

Our cost function will rely on this simplified cost model:

- Whenever an engine **operates** for a time step, we gain a **profit** of 1 unit
- A failure costs C units (i.e. the equivalent of C operation days)
- We never trigger maintenance before s time steps (safe interval)

Some comments:

- C is actually an offset over the cost of maintenance
- The last rule mimics using preventive maintenance as a fail-safe mechanism



The Cost Function

Normally, we would determine s and C by talking to a domain expert

...In our case we will pick reasonable values based on our data

- First, we collect all failure times:

```
In [3]: tr_failtimes = tr.groupby('machine')['cycle'].max()
```

- Then, we define s and C based on statistics:

```
In [4]: safe_interval = tr_failtimes.min()  
maintenance_cost = tr_failtimes.max()
```

- For the safe interval s , we choose the minimum failure time
- For the maintenance cost C we choose the largest failure time

We are talking about jet engines, so failing is BAD



Solving the Calibration Problem


We can **sample a range** of values for the θ parameter

...Then simply pick the value with the smallest cost

- The code in `optimize_threshold` can also plot the corresponding cost surface

```
In [5]: cmodel = util.RULCostModel(maintenance_cost, safe_interval)
th_range = np.linspace(0, 40, 100)
tr_thr = util.optimize_threshold(tr['machine'].values, tr_pred, th_range, cmodel, plot=True)
print(f'Optimal threshold for the training set: {tr_thr:.2f}')
```

Optimal threshold for the training set: 11.31



Evaluation

Finally, we can check how we are doing on the test set:

```
In [6]: tr_c, tr_f, tr_sl = cmodel.cost(tr['machine'].values, tr_pred, tr_thr, return_margin=True)
        ts_c, ts_f, ts_sl = cmodel.cost(ts['machine'].values, ts_pred, tr_thr, return_margin=True)
        print(f'Cost: {tr_c/len(tr_mcn):.2f} (training), {ts_c/len(ts_mcn):.2f} (test)')
```

```
Cost: -90.81 (training), -104.38 (test)
```

We can also evaluate the margin for improvement:

```
In [7]: print(f'Avg. fails: {tr_f/len(tr_mcn):.2f} (training), {ts_f/len(ts_mcn):.2f} (test)')
        print(f'Avg. slack: {tr_sl/len(tr_mcn):.2f} (training), {ts_sl/len(ts_mcn):.2f} (test)')
```

```
Avg. fails: 0.01 (training), 0.00 (test)
Avg. slack: 22.35 (training), 19.56 (test)
```

- Slack = distance between when we stop and the failure
- The results are actually quite good and we also generalize fairly well



Some Considerations

In principle, RUL regression is a very hard problem

- Our linearly decreasing RUL assumption is just a rough oversimplification
- ...RUL is inherently subject to stochasticity
- ...And depends on the how the machine **will be** used

But we **don't care, since RUL prediction was **not our true problem****

The real problem involved both **prediction and optimization**

- We had to optimize the NN parameters (to obtain good predictions)
- We had to optimize the threshold

The ultimate goal was to **reduce maintenance cost**

It's worth to keep in mind **the big picture**

- In a "predict, then optimize" setting
- ...Quality should be judged on the final cost

